Xuanming Zhang

Researcher in Humanoid Intelligence

+1 (608) 259-0539 zhangxm@stanford.edu web.stanford.edu/people/zhangxm XMZhangAI

Summary

Senior undergraduate at UW-Madison (ISci, GPA: 3.9/4.0) with research experience at Amazon, Stanford, Tsinghua, and Peking University. Visiting B.S. researcher at Stanford University (A+). My research focuses on Humanoid Intelligence: Cognitive Learning, World Interaction, and Policy Simulation. Published 10+ papers at top venues (NeurIPS Spotlight, ACL) with 300+ citations and h-index of 7.

Education

Stanford University May 2025 – Present
Visiting Researcher, Computer Science *Stanford, CA*
Grade: A+ • Finished Graduate Level STATS217, 200, EE364

University of Wisconsin-Madison Dec 2024 – May 2026
B.S. in Information Science (Transfer) *Madison, WI*
GPA: 3.9/4.0 • Honor

Peking University Sep 2022 – Dec 2024
Artificial Intelligence (EECS, Withdrawn) *Beijing, China*
National Elite Institute of Engineering Program

Research Experience

Amazon - Trustworthy AI Research May 2025 – Sep 2025
Research Scholar *Remote*

- Working with Dr. Rahul Gupta and Dr. Ousmane Amadou Dia on AI world interaction
- Developing evaluation frameworks for LLM deceptive behavior in workspace scenarios

UW-Madison - Deep Learning Lab Feb 2025 – Present
Research Assistant (Advisor: Prof. Sharon Li) *Madison, WI*

- Leading 3 research projects on LLM reasoning vulnerability and cognitive architectures
- Developed MetaMind for AI social intelligence through cognitive learning frameworks (NeurIPS 2025)

ByteDance - Flow Sep 2024 – Feb 2025
AI Architect Intern (Advisors: Xingnan Wang, Xuebin Lin, Muyang Li) *Beijing, China*

- Research solutions and deployment of Applet
- Co-Design Fourier Analysis Networks for LLM periodic pattern

Tsinghua University - CoAI Lab & Zhipu-AI Apr 2024 – Feb 2025
Research Scientist Intern (Advisors: Prof. Minlie Huang, Prof. Hongning Wang) *Beijing, China*

- Led Seeker to approach robustness code generation for LLMs (TOSEM Under Review)
- Co-Led SocialEval to evaluate social intelligence of LLMs (ACL 2025)

Peking University - Key Lab of High Confidence Software Technologies Jun 2023 – May 2024
Research Assistant (Advisors: Prof. Zhi Jin, Prof. Ge Li) *Beijing, China*

- Co-Led EvoCodeBench evaluating code generation for LLMs in realistic development (NeurIPS 2024)
- Co-Led MGI to evaluate, mitigate data contamination for LLMs (ACL 2024)

Zen-AI Jan 2023 – Mar 2023
SDE Intern (Advisor: Libing Yan) *Beijing, China*

Selected Publications

12. Cognition-of-Thought Elicits Social-Aligned Reasoning in Large Language Models — **ICLR 2026 (Review)**
X. Zhang, Y. Chen, M. Yeh, S. Li

11. Simulating and Understanding Deceptive Behaviors in Long-Horizon Interactions — **ICLR 2026 (Review)**
Y. Xu*, **X. Zhang***, M. Yeh, ..., S. Li

10. MetaMind: Modeling Human Social Thoughts with Metacognitive Multi-Agent Systems — **NeurIPS 2025 Spotlight**
X. Zhang, Y. Chen, M. Yeh, S. Li

9. Theoretical Proof that Generated Text in the Corpus Leads to the Collapse of Auto-regressive Language Models — **Sci China Inf Sci (Review)**
L. Wang, X. Shi, ..., **X. Zhang**, ..., H. Mei

8. Human Decision-making is Susceptible to AI-driven Manipulation — **Nature Comm. (Revision)**
S. Sabour, J. Liu, ..., **X. Zhang**, ..., M. Huang

7. SocialEval: Evaluating Social Intelligence of Large Language Models — **ACL 2025**
J. Zhou*, Y. Chen*, Y. Shi*, **X. Zhang***, ..., M. Huang (*Equal contribution)

6. Towards Exception Safety Code Generation with Intermediate Representation Agents Framework — **TOSEM (Review)**
X. Zhang, Y. Chen, Y. Yuan, M. Huang

5. EvoCodeBench: An Evolving Code Generation Benchmark with Domain-Specific Evaluations — **NeurIPS 2024**
J. Li, G. Li, **X. Zhang**, Y. Dong, Z. Jin

4. Generalization or Memorization: Data Contamination and Trustworthy Evaluation for Large Language Models — **TPAMI (Review)**
Y. Dong, X. Jiang, **X. Zhang**, ..., M. Yang, G. Li

3. DevEval: A Manually-Annotated Code Generation Benchmark Aligned with Real-World Code Repositories — **ACL 2024**
J. Li, G. Li, ..., **X. Zhang**, ..., F. Huang, Y. Li

2. Deep Graph Learning for Industrial Carbon Emission Analysis and Policy Impact — **NeurIPS 2025 AI for Science Workshop**
X. Zhang

1. The average human lifespan is only 18 years — **Beijing Evening News**
X. Zhang

Key Projects

RAG4Code — PKU 2024
Achieved SOTA on CrossCodeEval with dependency-chain-retrieval for cross-file code completion

Common Exception Enumeration (CEE) — THU 2024
Curated knowledge base systematizing exception handling patterns across programming languages

Applet - Mobile UI Agent — ByteDance 2024
Lightweight intent-driven framework using Swift/CoreML for real-time multimodal interaction

H3O - Startup (CTO) — 2022
Received 500K RMB investment from XbotPark; led architecture and prototype development

Technical Skills

Languages: Python, C++, Swift, Java, Shell

ML/AI: Transformers, PEFT, vLLM, CoreML

Tools: Git, Docker, LaTeX, Distributed Systems

Selected Awards

National Innovation Excellence Award

Chinese Mathematics Competition First Prize

Outstanding Student President Scholarship

Service & Leadership

Academic: Program Committee Member for ICLR, NeurIPS, OOPSLA, AAAI • Student Member at CCF, IEEE

Leadership: College Student Union President (BNDS 2020, EIE 2022)

Certifications: National Level II Psychological Counselor (PCIA)